

(a) Load 104

| Step | RTN | PC | IR | MAR | MBR | AC |
|------------------|---------------------------|-----|-------|-------|-------|-------|
| (initial values) | | 100 | ----- | ----- | ----- | ----- |
| Fetch | MAR \leftarrow PC | 100 | ----- | 100 | ----- | ----- |
| | IR \leftarrow M [MAR] | 100 | 1104 | 100 | ----- | ----- |
| | PC \leftarrow PC + 1 | 101 | 1104 | 100 | ----- | ----- |
| Decode | MAR \leftarrow IR[11-0] | 101 | 1104 | 104 | ----- | ----- |
| | (Decode IR[15-12]) | 101 | 1104 | 104 | ----- | ----- |
| Get operand | MBR \leftarrow M [MAR] | 101 | 1104 | 104 | 0023 | ----- |
| Execute | AC \leftarrow MBR | 101 | 1104 | 104 | 0023 | 0023 |

(b) Add 105

| Step | RTN | PC | IR | MAR | MBR | AC |
|------------------|---------------------------|-----|------|-----|------|------|
| (initial values) | | 101 | 1104 | 104 | 0023 | 0023 |
| Fetch | MAR \leftarrow PC | 101 | 1104 | 101 | 0023 | 0023 |
| | IR \leftarrow M [MAR] | 101 | 3105 | 101 | 0023 | 0023 |
| | PC \leftarrow PC + 1 | 102 | 3105 | 101 | 0023 | 0023 |
| Decode | MAR \leftarrow IR[11-0] | 102 | 3105 | 105 | 0023 | 0023 |
| | (Decode IR[15-12]) | 102 | 3105 | 105 | 0023 | 0023 |
| Get operand | MBR \leftarrow M [MAR] | 102 | 3105 | 105 | FFE9 | 0023 |
| Execute | AC \leftarrow AC + MBR | 102 | 3105 | 105 | FFE9 | 000C |

(c) Store 106

| Step | RTN | PC | IR | MAR | MBR | AC |
|------------------|---------------------------|-----|------|-----|------|------|
| (initial values) | | 102 | 3105 | 105 | FFE9 | 000C |
| Fetch | MAR \leftarrow PC | 102 | 3105 | 102 | FFE9 | 000C |
| | IR \leftarrow M [MAR] | 102 | 2106 | 102 | FFE9 | 000C |
| | PC \leftarrow PC + 1 | 103 | 2106 | 102 | FFE9 | 000C |
| Decode | MAR \leftarrow IR[11-0] | 103 | 2106 | 106 | FFE9 | 000C |
| | (Decode IR[15-12]) | 103 | 2106 | 106 | FFE9 | 000C |
| Get operand | (not necessary) | 103 | 2106 | 106 | FFE9 | 000C |
| Execute | MBR \leftarrow AC | 103 | 2106 | 106 | 000C | 000C |
| | M [MAR] \leftarrow MBR | 103 | 2106 | 106 | 000C | 000C |

FIGURE 4.14 A Trace of the Program to Add Two Numbers

4.11 A DISCUSSION ON ASSEMBLERS

In the program shown in Table 4.3, it is a simple matter to convert from the assembly language instruction Load 104, for example, to the machine language instruction 1104 (in hex). But why bother with this conversion? Why not just write in machine code? Although it is very efficient for computers to see these instructions as binary numbers, it is difficult for human beings to understand and program in sequences of 0s and 1s. We prefer words and symbols over long numbers, so it seems a natural solution to devise a program that does this simple conversion for us. This program is called an **assembler**.

4.11.1 What Do Assemblers Do?

An assembler's job is to convert assembly language (using mnemonics) into machine language (which consists entirely of binary values, or strings of 0s and 1s). Assemblers take a programmer's assembly language program, which is really a symbolic representation of the binary numbers, and convert it into binary instructions, or the machine code equivalent. The assembler reads a **source file** (assembly program) and produces an **object file** (the machine code).

Substituting simple alphanumeric names for the opcodes makes programming much easier. We can also substitute **labels** (simple names) to identify or name particular memory addresses, making the task of writing assembly programs even simpler. For example, in our program to add two numbers, we can use labels to indicate the memory addresses, thus making it unnecessary to know the exact memory address of the operands for instructions. Table 4.4 illustrates this concept.

When the address field of an instruction is a label instead of an actual physical address, the assembler still must translate it into a real, physical address in main memory. Most assembly languages allow for labels. Assemblers typically specify formatting rules for their instructions, including those with labels. For example, a label might be limited to three characters and may also be required to occur as the first field in the instruction. MARIE requires labels to be followed by a comma.

Labels are nice for programmers. However, they make more work for the assembler. It must make two passes through a program to do the translation. This

| Address | Instruction | |
|---------|-------------|---|
| 100 | Load | X |
| 101 | Add | Y |
| 102 | Store | Z |
| 103 | Halt | |
| 104 X, | 0023 | |
| 105 Y, | FFE9 | |
| 106 Z, | 0000 | |

TABLE 4.4 An Example Using Labels

means the assembler reads the program twice, from top to bottom each time. On the first pass, the assembler builds a set of correspondences called a **symbol table**. For the above example, it builds a table with three symbols: X, Y, and Z. Because an assembler goes through the code from top to bottom, it cannot translate the entire assembly language instruction into machine code in one pass; it does not know where the data portion of the instruction is located if it is given only a label. But after it has built the symbol table, it can make a second pass and “fill in the blanks.”

In the above program, the first pass of the assembler creates the following symbol table:

| | |
|---|-----|
| X | 104 |
| Y | 105 |
| Z | 106 |

It also begins to translate the instructions. After the first pass, the translated instructions would be incomplete as follows:

| | | | |
|---|---|---|---|
| 1 | X | | |
| 3 | Y | | |
| 2 | Z | | |
| 7 | 0 | 0 | 0 |

On the second pass, the assembler uses the symbol table to fill in the addresses and create the corresponding machine language instructions. Thus, on the second pass, it would know that X is located at address 104, and would then substitute 104 for the X. A similar procedure would replace the Y and Z, resulting in:

| | | | |
|---|---|---|---|
| 1 | 1 | 0 | 4 |
| 3 | 1 | 0 | 5 |
| 2 | 1 | 0 | 6 |
| 7 | 0 | 0 | 0 |

Because most people are uncomfortable reading hexadecimal, most assembly languages allow the data values stored in memory to be specified as binary, hexadecimal, or decimal. Typically, some sort of **assembler directive** (an instruction specifically for the assembler that is not supposed to be translated into machine code) is given to the assembler to specify which base is to be used to interpret the value. We use DEC for decimal and HEX for hexadecimal in MARIE’s assembly language. For example, we rewrite the program in Table 4.4 as shown in Table 4.5.

Instead of requiring the actual binary data value (written in HEX), we specify a decimal value by using the directive DEC. The assembler recognizes this directive and converts the value accordingly before storing it in memory. Again, directives are not converted to machine language; they simply instruct the assembler in some way.

Another kind of directive common to virtually every programming language is the **comment delimiter**. Comment delimiters are special characters that tell the

| Address | Instruction | |
|---------|-------------|------|
| 100 | Load | X |
| 101 | Add | Y |
| 102 | Store | Z |
| 103 | Halt | |
| 104 X, | DEC | 35 |
| 105 Y, | DEC | -23 |
| 106 Z, | HEX | 0000 |

TABLE 4.5 An Example Using Directives for Constants

assembler (or compiler) to ignore all text following the special character. MARIE's comment delimiter is a front slash ("/"), which causes all text between the delimiter and the end of the line to be ignored.

4.11.2 Why Use Assembly Language?

Our main objective in presenting MARIE's assembly language is to give you an idea of how the language relates to the architecture. Understanding how to program in assembly goes a long way toward understanding the architecture (and vice versa). Not only do you learn basic computer architecture, but you also can learn exactly how the processor works and gain significant insight into the particular architecture on which you are programming. There are many other situations where assembly programming is useful.

Most programmers agree that 10% of the code in a program uses approximately 90% of the CPU time. In time-critical applications, we often need to optimize this 10% of the code. Typically, the compiler handles this optimization for us. The compiler takes a high-level language (such as C++) and converts it into assembly language (which is then converted into machine code). Compilers have been around a long time and in most cases they do a great job. Occasionally, however, programmers must bypass some of the restrictions found in high-level languages and manipulate the assembly code themselves. By doing this, programmers can make the program more efficient in terms of time (and space). This hybrid approach (most of the program written in a high-level language, with part rewritten in assembly) allows the programmer to take advantage of the best of both worlds.

Are there situations in which entire programs should be written in assembly language? If the overall size of the program or response time is critical, assembly language often becomes the language of choice. This is because compilers tend to obscure information about the cost (in time) of various operations and programmers often find it difficult to judge exactly how their compiled programs will perform. Assembly language puts the programmer closer to the architecture and, thus, in firmer control. Assembly language might actually be necessary if the programmer wishes to accomplish certain operations not available in a high-level language.

A perfect example, in terms of both response performance and space-critical design, is found in **embedded systems**. These are systems in which the computer is integrated into a device that is typically not a computer. Embedded systems

must be reactive and often are found in time-constrained environments. These systems are designed to perform either a single instruction or a very specific set of instructions. Chances are you use some type of embedded system every day. Consumer electronics (such as cameras, camcorders, cellular phones, PDAs, and interactive games), consumer products (such as washers, microwave ovens, and washing machines), automobiles (particularly engine control and antilock brakes), medical instruments (such as CAT scanners and heart monitors), and industry (for process controllers and avionics) are just a few of the examples of where we find embedded systems.

The software for an embedded system is critical. An embedded software program must perform within very specific response parameters and is limited in the amount of space it can consume. These are perfect applications for assembly language programming. We delve deeper into this topic in Chapter 10.

4.12 EXTENDING OUR INSTRUCTION SET

Even though MARIE's instruction set is sufficient to write any program we wish, there are a few instructions we can add to make programming much simpler. We have 4 bits allocated to the opcode, which implies we can have 16 unique instructions, and we are using only 9 of them. We add the instructions from Table 4.6 to extend our instruction set.

The JnS (jump-and-store) instruction allows us to store a pointer to a return instruction and then proceeds to set the PC to a different instruction. This enables us to call procedures and other subroutines, and then return to the calling point in our code once the subroutine has finished. The Clear instruction moves all 0s into the accumulator. This saves the machine cycles that would otherwise be expended in loading a 0 operand from memory.

The AddI instruction (as well as the JumpI instruction) uses a different **addressing mode**. All previous instructions assume the value in the data portion of the instruction is the **direct address** of the operand required for the instruction. The AddI instruction uses the **indirect addressing mode** (we present more on addressing modes in Chapter 5). Instead of using the value found at location X as the actual address, we use the value found in X as a pointer to a new memory

| Instruction Number (hex) | Instruction | Meaning |
|--------------------------|-------------|---|
| 0 | JnS X | Store the PC at address X and jump to X + 1. |
| A | Clear | Put all zeros in AC. |
| B | AddI X | Add indirect: Go to address X. Use the value at X as the actual address of the data operand to add to AC. |
| C | JumpI X | Jump indirect: Go to address X. Use the value at X as the actual address of the location to jump to. |

TABLE 4.6 MARIE's Extended Instruction Set

location that contains the data we wish to use in the instruction. For example, if we have the instruction AddI 400, we would go to location 400, and assuming we found the value 240 stored at location 400, we would go to location 240 to get the actual operand for the instruction. We have, essentially, allowed for pointers in our language.

Returning to our discussion of register transfer notation, our new instructions are represented as follows:

JnS

```
MBR ← PC
MAR ← X
M[MAR] ← MBR
MBR ← X
AC ← 1
AC ← AC + MBR
PC ← AC
```

Clear

```
AC ← 0
```

AddI X

```
MAR ← X
MBR ← M[MAR]
MAR ← MBR
MBR ← M[MAR]
AC ← AC + MBR
```

JumpI X

```
MAR ← X
MBR ← M[MAR]
PC ← MBR
```

Table 4.7 summarizes MARIE's entire instruction set.

Let's look at some examples using the full instruction set.

☰ EXAMPLE 4.1 Here is an example using a loop to add five numbers:

| Address | Instruction | | |
|---------|-------------|------|---|
| 100 | Load | Addr | /Load address of first number to be added |
| 101 | Store | Next | /Store this address as our Next pointer |
| 102 | Load | Num | /Load the number of items to be added |
| 103 | Subt | One | /Decrement |

| | | | |
|-----|----------|-------|---|
| 104 | Store | Ctr | /Store this value in Ctr to control looping |
| 105 | Loop, | Load | Sum /Load the Sum into AC |
| 106 | | AddI | Next /Add the value pointed to by location Next |
| 107 | | Store | Sum /Store this sum |
| 108 | | Load | Next /Load Next |
| 109 | | Add | One /Increment by one to point to next address |
| 10A | | Store | Next /Store in our pointer Next |
| 10B | | Load | Ctr /Load the loop control variable |
| 10C | | Subt | One /Subtract one from the loop control variable |
| 10D | | Store | Ctr /Store this new value in loop control variable |
| 10E | Skipcond | 000 | /If control variable < 0, skip next /instruction |
| 10F | Jump | Loop | /Otherwise, go to Loop |
| 110 | | Halt | /Terminate program |
| 111 | Addr, | Hex | 117 /Numbers to be summed start at location 118 |
| 112 | Next, | Hex | 0 /A pointer to the next number to add |
| 113 | Num, | Dec | 5 /The number of values to add |
| 114 | Sum, | Dec | 0 /The sum |
| 115 | Ctr, | Hex | 0 /The loop control variable |
| 116 | One, | Dec | 1 /Used to increment and decrement by 1 |
| 117 | | Dec | 10 /The values to be added together |
| 118 | | Dec | 15 |
| 119 | | Dec | 20 |
| 11A | | Dec | 25 |
| 11B | | Dec | 30 |

Note: Line numbers in program are given for information only and are not required for use in the MarieSim environment.

Although the comments are reasonably explanatory, let's walk through Example 4.1. Recall that the symbol table stores [label, location] pairs. The Load Addr instruction becomes Load 111, because Addr is located at physical memory address 111. The value of 117 (the value stored at Addr) is then stored in Next. This is the pointer that allows us to "step through" the five values we are adding (located at addresses 117, 118, 119, 11A, and 11B). The Ctr variable keeps track of how many iterations of the loop we have performed. Because we are checking to see if Ctr is negative to terminate the loop, we start by subtracting one from Ctr. Sum (with an initial value of 0) is then loaded in the AC. The loop begins, using Next as the address of the data we wish to add to the AC. The Skipcond statement terminates the loop when Ctr is negative by skipping the unconditional branch to the top of the loop. The program then terminates when the Halt statement is executed.

Example 4.2 shows how you can use the Skipcond and Jump instructions to perform selection. Although this example illustrates an if/else construct, you can easily modify this to perform an if/then structure, or even a case (or switch) structure.